

C++20 STL Features: 1 Year of Development on GitHub

Stephan T. Lavavej "Steh-fin Lah-wah-wade"


Principal Software Engineer, Visual C++ Libraries

stl@microsoft.com

@StephanTLavavej

Getting Started

- Please hold your questions until the end
 - Write down the slide numbers
- Part 0: Overview
 - What's happened in the last year
- Part 1: C++20 STL Features
 - Everything here is Standard, except as noted
- Part 2: GitHub Development
 - For contributors and observers




Overview

Part 0

CppCon 2019 -> CppCon 2020

- Announced at CppCon 2019
 - github.com/microsoft/STL
 - Apache License v2.0 with LLVM Exception
- Implemented ~50 C++20 features
 - Majority from our amazing contributors
 - Extensively reviewed and tested
- Released VS 2019 16.5, 16.6, 16.7, 16.8 Preview 3
- GitHub migration ongoing
 - Build system (native desktop), test suite, issue database
- Goal: Complete C++20 in 2020
 - No promises, but we're working hard; refers to repo only

C++20 Features: 60 -> 23 Left



C++20 STL Features

Part 1

Comparing Integers

Usual Arithmetic Conversions

```
#include <iostream>
using namespace std;
int main() {
    cout << boolalpha;
    short s          = -1;
    unsigned int ui = 1729;
    cout << (s < ui) << endl;
}
```

- What does this print?
 - warning C4018: '<': signed/unsigned mismatch
 - warning: comparison of integers of different signs: 'short' and 'unsigned int' [-Wsign-compare]
 - This prints **false** 🙈

Integer Comparison Functions

- [GH-621](#) implemented by Neargye

```
#include <iostream>
#include <utility>
using namespace std;
int main() {
    cout << boolalpha;
    short s          = -1;
    unsigned int ui = 1729;
    cout << cmp_less(s, ui) << endl;
}
```

- This prints **true** ✓

<utility>

```
template <class T, class U>
    constexpr bool cmp_equal(T t, U u) noexcept;
template <class T, class U>
    constexpr bool cmp_not_equal(T t, U u) noexcept;
template <class T, class U>
    constexpr bool cmp_less(T t, U u) noexcept;
template <class T, class U>
    constexpr bool cmp_greater(T t, U u) noexcept;
template <class T, class U>
    constexpr bool cmp_less_equal(T t, U u) noexcept;
template <class T, class U>
    constexpr bool cmp_greater_equal(T t, U u) noexcept;
template <class R, class T>
    constexpr bool in_range(T t) noexcept;
```

constexpr Algorithms

C++20 constexpr Everything

- constexpr enables compile-time programming without arcane templates
- Algorithms:
 - P0202R3 constexpr <algorithm> And exchange() ([GH-425](#))
 - P0879R0 constexpr Algorithms, Part II ([GH-425](#), BillyONeal)
 - P1645R1 constexpr <numeric> Algorithms ([GH-399](#), Neargye)
- More:
 - P0415R1 constexpr <complex> Again ([GH-367](#), Neargye)
 - P0595R2 is_constant_evaluated() ([GH-353](#), Jennifer Yao – MSVC compiler)
 - P1006R1 constexpr pointer_to() ([GH-397](#), AdamBucior)
 - P1023R0 constexpr array Comparisons ([GH-599](#), Weheineman)
 - P1032R1 Miscellaneous constexpr ([GH-491](#), miscoco)
 - P1065R2 constexpr INVOKE ([GH-703](#), AdamBucior)
- Work in progress (mnatsuhara):
 - [GH-37](#) P0784R7 Library Support For More constexpr Containers
 - [GH-43](#) P0980R1 constexpr string
 - [GH-45](#) P1004R2 constexpr vector

constexpr In Everyday Programming

- constexpr has steadily increased in power
 - It's just code that runs on your machine instead of the user's
- Example: lookup tables
 - Everyone uses lookup tables
 - Arrays are maximally space-efficient, very time-efficient
 - Sorted arrays: `binary_search`, `lower_bound`, `equal_range`
- Now you can use constexpr algorithms
 - Easy: `static_assert` with `constexpr is_sorted`
 - `string_view` is also constexpr; not limited to numeric data
 - More advanced: constexpr sort your lookup tables

static_assert is_sorted

- Compile-time evaluation = no run-time codegen, even in debug builds

```
#include <algorithm>
#include <array>
#include <filesystem>
#include <string>
#include <string_view>
using namespace std;
int main() {
    static constexpr array skipped_extensions{".dll"sv, ".exe"sv, ".obj"sv};
    static_assert(ranges::is_sorted(skipped_extensions));
    for (const auto& ent : filesystem::recursive_directory_iterator{"."}) {
        const string extension = ent.path().extension().string();
        if (ranges::binary_search(skipped_extensions, extension)) {
            continue;
        }
        // ...
    }
}
```

Erasing Elements

Container Erasure Strategies

- Erasing unwanted elements before C++20:
 - vector-like: erase-remove idiom
 - list-like: remove/remove_if member functions
 - map-like: handwritten loop calling m.erase(iter)
- Many potential hazards 🙈
 - `remove_if(v.begin(), v.end(), pred);` 🐞
 - `v.erase(remove_if(v.begin(), v.end(), pred));` 🐞
 - Quadratic complexity `vec.erase(iter)` loop 🐞
 - Invalidating iterators while looping 🐞
 - Skipping elements while looping 🐞

Uniform Container Erasure

- [GH-236](#) and [GH-566](#) implemented by SuperWig

```
#include <iostream>
#include <string>
#include <vector>
using namespace std;
void print(const vector<string>& v) {
    for (const auto& e : v) { cout << e << " "; }
    cout << "\n";
}
int main() {
    vector<string> v{"bear", "dog", "cat", "lion", "ox", "dog"};
    auto pred = [](const auto& s) { return s.size() > 3; };
    erase_if(v, pred);
    print(v); // dog cat ox dog
    erase(v, "dog");
    print(v); // cat ox
}
```

atomic_ref

atomic vs. atomic_ref

- `atomic<T>`: atomic access and storage for T
- `atomic_ref<T>`: atomic access for separate T
- Scenarios (see [P0019R8](#)):
 - Distinct single-threaded and multi-threaded parts
 - Pay the cost for atomic operations only when necessary
 - Distinct non-conflicting and conflicting parts
 - Layout compatibility; works with T within C structs
- Near-identical interface
 - Encapsulates platform/architecture-specific code
 - `atomic/atomic_ref` support `wait/notify_one/notify_all`

Highly Artificial Example

- [GH-843](#) implemented by AlexGuteniev

```
#include <atomic>
#include <iostream>
#include <thread>
using namespace std;
int main() {
    int i{500};
    i += 500; // ordinary read/write
    { atomic_ref atom{i};
        thread t1{[&atom] { for (int val{0}, x{0}; x < 70;) {
            if (atom.compare_exchange_weak(val, val + 10)) { ++x; }}}};
        thread t2{[&atom] { for (int val{0}, y{0}; y < 29;) {
            if (atom.compare_exchange_weak(val, val + 1)) { ++y; }}}};
        t1.join(); t2.join(); }
    cout << i << endl; // ordinary read, 1729
}
```

span

span Encapsulates Ptr-Len Pairs

- span is a non-owning view of contiguous elements
- `span<T>` is like a pair of `T*` and `size_t`, but better
 - spans are implicitly constructible from arrays and vectors
 - Avoids ownership confusion, accidental pointer arithmetic
 - `span<T> → span<const T>`; `span<Derived> ✗ span<Base>`
- Shallow const, like `T*`
 - `span<const T>` can be reassigned, can't modify elements
 - `const span<T>` can't be reassigned, can modify elements
- MSVC's `std::span` checks debug; `gs1::span` always

Contiguous Range Constructor

- [GH-142](#), [GH-500](#), [GH-587](#) implemented by miscco

```
#include <array>
#include <iostream>
#include <span>
#include <vector>
using namespace std;

void print(const span<const int> s) { // not a template!
    for (const auto& e : s) { cout << e << " "; }
    cout << endl;
}

int main() {
    static constexpr int classic[]{1, 2, 3};
    print(classic);
    print(array{4, 5, 6});
    print(vector{7, 8, 9});
}
```

Contiguous > Random-Access

- Can we `print(deque{-1, -1, -1});` ?
 - candidate function not viable: **no known conversion** from '`std::deque<int, std::allocator<int> >`' to '`const span<const int>`' for 1st argument
- This constructor is powered by C++20 ranges!
 - Specifically `ranges::contiguous_range` (and more)
 - Extensible to user-defined contiguous ranges

And More!

So Many New Features, Including:

- `atomic<shared_ptr<T>>`, `atomic<weak_ptr<T>>`
 - [GH-601](#) by AdamBucior
- `<bit>` (`bit_cast`, `rotating/counting`, `power-of-2`)
 - Several PRs by barcharcraz
- `make_shared()` For Arrays
 - [GH-309](#) by AdamBucior, tested by Weheineman
- `midpoint()`, `lerp()`
 - Paolo Torres and BillyONeal initially implemented this
 - [GH-1048](#) by statementreply completed this
- `<numbers>` Math Constants
 - [GH-261](#) by SuperWig

GitHub Development

Part 2

How We Use GitHub

- Code: linear history, few feature branches
- Issues: cxx20, LWG, bug, performance, etc.
- Pull Requests: Used for all development
- Continuous Integration: Azure Pipelines
 - Code format validation: clang-format, line length, etc.
- Projects: Code reviews, C++20 features, etc.
- Wiki: Changelog, checklists, other resources
- Status Chart: Generated via GitHub's GraphQL API


Code: github.com/microsoft/STL

- Product: `stl/inc`, `stl/src`
- Build system: CMake/Ninja
 - `vcpkg` submodule acquires Boost.Math for Special Math
- `tests/std`, `tests/libcxx`, `tests/tr1` (legacy)
 - `llvm-project` submodule for `libc++`'s test suite
 - Uses Python and `lit` (LLVM Integrated Tester)
- Linear history, squashed commits
- Feature branches are uncommon
 - `feature/spaceship`, `feature/format`

Issues

- **cxx20** issues track all remaining C++20 features
 - [GH-39](#) P0896R4 <ranges>
- **LWG** issues track bugfixes in the Standard itself
 - [GH-333](#) LWG-3070 path::lexically_relative
- **bug, performance, throughput, enhancement**
 - Porting from Developer Community, Azure Boards, todos
 - [GH-713](#) boyer_moore_searcher
- **vNext** issues affect binary compatibility
 - Will start vNext after finishing C++20

cxx20 Issues


 microsoft/STL

[Unwatch releases](#) 227 [Unstar](#) 5.9k [Fork](#) 691


[Code](#) [Issues 322](#) [Pull requests 28](#) [Discussions](#) [Projects 7](#) [Wiki](#) [...](#)


P0980R1 constexpr std::string #43


[Open](#) StephanTLavavej opened this issue on Sep 5, 2019 · 0 comments


 StephanTLavavej commented on Sep 5, 2019 • edited [Member](#) [...](#)

`P0980R1 constexpr std::string`
Feature-test macro as of WG21-N4842:
`#define __cpp_lib_constexpr_string 201907L`

 StephanTLavavej added the `cxx20` label on Sep 5, 2019

 mnatsuhara self-assigned this on Apr 28

 StephanTLavavej added the `work in progress` label on Apr 29

 StephanTLavavej mentioned this issue on Apr 29

[P0784R7 Library Support For More constexpr Containers #37](#) [Open](#)

Successfully merging a pull request may close this issue.

Assignees [mnatsuhara](#)

Labels `cxx20` `work in progress`

Projects `C++20 Features` Investigating

Milestone

Conformance

Linked pull requests

Pull Requests

- [GH-724](#) Fix boyer_moore_searcher
 - Fixed a 43-year-old bug by implementing a 40-year-old fix
 - Explained bug, fix, ABI, testing; others provided more info
- [GH-142](#) P0122R7 ``
 - 691 comments!
- [GH-1173](#) Fix ostr `<< flt` precision
 - Fixed long-standing bugs while preserving ABI
- Every PR is extensively reviewed by 2 maintainers
 - Keeps code at production quality: always ready to ship
 - Helps the team understand the code and related features

Code Review Philosophy

- The Standard Library has:
 - Unusually well-specified preconditions and postconditions
 - Unusually strict correctness and performance requirements
 - Indefinite lifespan: old codebase, yet must keep evolving
- Careful code review avoids regressions, new bugs
 - Especially important for binary compatibility
- Codebase consistency is important, especially now
 - Ideally, looks like the work of a single author
 - Consistent code -> consistent behavior, fewer unique bugs
 - Makes inconsistent code stand out as unusual or incorrect


Code Reviews Project

The screenshot displays the 'Code Reviews Project' dashboard, which is organized into four main columns representing different stages of the review process. The top navigation bar includes the GitHub logo, the path 'microsoft / STL > Projects > Code Reviews', a search bar for 'Filter cards', and buttons for 'Add cards', 'Exit fullscreen', and 'Menu'.

- Work In Progress (13 items):** This column contains three pull requests. The first, 'Smart Pointer Creation With Default Initialization' (#778) by Weheineman, is marked 'cx20' and has 'Changes requested'. The second, '<tuple>: tuple_cat for tuple-like types' (#461) by ArtemSarmini, is marked 'enhancement' and has 'Changes requested'. The third, '<memory>(mostly): Uses-Allocator and guaranteed copy elision For piecewise construction' (#421) by ArtemSarmini, is marked 'cx20' and has 'Review required'.
- Initial Review (12 items):** This column contains three pull requests. The first, '<chrono> Partially implement P0355R7' (#323) by SuperWig, is marked 'cx20' and has 'Changes requested'. The second, 'Efficient Access To basic_stringbuf's Buffer' (#919) by Berrysoft, is marked 'cx20' and has 'Changes requested'. The third, '[xutility] Modernize_Ptr_meow_cat to use variable templates' (#872) by miscco, is marked 'throughput' and has 'Changes requested'.
- Final Review (3 items):** This column contains two pull requests. The first, 'Revert "Revert "Implement common_iterator (#1092)" (#1265)"' (#1268) by CaseyCarter, is marked 'cx20', 'ranges', and 'uncharted', and has 'Review required'. The second, 'Integrate P1391 into views::take and views::drop' (#1271) by CaseyCarter, is marked 'cx20' and 'ranges', and has 'Review required'.
- Ready To Merge (2 items):** This column contains one pull request, '<yvals_core.h>: Updated _MSVC_STL_UPDATE to September 2020' (#1272) by Madhvan28, marked 'enhancement', with 'Changes approved'.

At the bottom of each column, there is an 'Automated as' status: 'In progress' for Work In Progress, 'To do' for Initial Review, and 'To do' for Final Review. A 'Manage' button is located at the bottom right of each column.

Status Chart: Monthly Merged PRs



Year 1 ends.
Year 2 begins!

More Info

- Links
 - Repository: github.com/microsoft/STL
 - Changelog: github.com/microsoft/STL/wiki/Changelog
 - Status Chart: microsoft.github.io/STL/
 - C++20: wg21.link/n4861
- Questions
 - GitHub Discussions tab
 - Discord server (see README)

Bonus Slides

Compile-Time/Run-Time Hybrid

```
template <class _Rx, class _Ty>
_NODISCARD constexpr bool in_range(const _Ty _Value) noexcept {
    // ... see <utility> lines 721-745 ...
    constexpr auto _Ty_min = _Min_limit<_Ty>();
    constexpr auto _Rx_min = _Min_limit<_Rx>();
    if constexpr (_STD cmp_less(_Ty_min, _Rx_min)) {
        if (_Value < _Ty{_Rx_min}) {
            return false;
        }
    }
    // ... similarly for _Value > _Ty{_Rx_max} ...
    return true;
}
```

Associative Erasure

- `std::erase_if()` is linear time
 - Inspects the entire `value_type`
- Associative containers have member `.erase(key)`
 - Inspects only the key, using the container's predicate
 - Ordered: "Logarithmic" time, $O(K + \log N)$
 - Unordered: "Constant" time, average $O(K)$, worst $O(N)$
- Only `std::erase_if()` is provided for associative
 - Avoids potential confusion

Continuous Integration

- Scripts prepare Azure Virtual Machine Scale Sets
 - Currently up to 12 VMs, each with 16 cores
 - VMs install VS (with Clang, CMake, Ninja), Python, CUDA
- We enforce clang-format for product/test code
 - Saves an incredible amount of time
 - Our parallelize tool runs it quickly, only on C++ files
 - Failed checks display the edits that clang-format wants
- Building all architectures is very fast
- Testing x86/x64 takes about an hour
 - Many tests, multiplied by many compiler configurations

Projects


- Code Reviews (primary dashboard!)
 - Work In Progress -> Initial Review -> Final Review -> Ready To Merge -> Done
- GitHub Migration
- C++20 Features
 - Available -> Investigating -> Reviewing PR -> Done
- Maintainer Priorities
- Large features: Chrono, Format, Spaceship

Wiki

- Changelog
 - Infeasible: Look at commit history between releases
 - Old process: Record commits, write them up later
 - New process: Merge PR, add Changelog entry immediately
- Advice for Reviewing Pull Requests
- MSVC STL Contributors Meeting Notes
- Macro `_MSVC_STL_UPDATE`
 - Updated every month by new contributors
- Custom Autolinks
 - DevCom-724444, LLVM-41915, LWG-3080, WG21-P1209

Status Chart

STL Status Chart



Status Chart: Pull Request Age

Pull Request Age

