

Design patterns for error handling in C++ programs using parallel algorithms and executors

Mark Hoemmen*

mhoemmen@stellarscience.com

CppCon 2020

* hoʊ'mən, or hæm'mən; he/him

Who am I?

- > 10 years post-PhD experience writing parallel C++ for science and engineering
- Background: Parallel algorithms for big linear algebra problems
- 1st WG21: Nov 2017
- Started new job at Stellar Science in March




Eschew raw pointers

Outline

- Parallelism makes error handling harder...
- ...C++ parallel algorithms and tasks specifically
- Message Passing Interface (MPI): 3 decades of distributed-memory parallel programming
- MPI teaches design patterns to detect and handle recoverable errors

What is “parallel”?

- **Use multiple hardware resources**
 - Nodes, cores, SIMD, ...
- **To accomplish >1 work item at the same time**
- **To improve performance**
 - Latency,
 - Throughput, or
 - Responsiveness




ILLIAC IV: the 1st massively parallel computer (image: Steve Jurvetson)

Parallel hinders error handling

- ... because parallelism relaxes execution order
 - Deliberately, to improve performance
- Errors interrupt execution; handling constrains order
- Errors could lead to deadlock (waiting forever)
 - e.g., 1 worker drops out before collective synchronization
- Correct handling requires communication
 - Data movement, or synchronization (same thing)
 - Stop other workers from waiting forever
 - Propagate and combine error info from workers

No zero-overhead solution

- Error handling requires communication
- Communication is expensive
- Making C++ “do it for you” won’t be free
- ➔ If you want a zero-overhead solution, ...
- ... you, the coder, will need to handle errors



(there is no free lunch)

Parallelism in Standard C++

- **Parallel (C++17) algorithms**
 - `for_each`, `reduce`, `transform`, `sort`, ...
 - `ExecutionPolicy`: Permitted changes in execution order
 - Throw in loop body → `terminate (*)`
- **Asynchronous tasks (C++11 `async`)**
 - Uncaught exception in task gets captured
 - Waiting on result throws passed-along exception
- **P0443 executors + P1897 asynchronous algorithms**
 - Separate path for handling ancestor task's uncaught exception
 - `when_all`: Express dependency on >1 tasks
 - If >1 parent throws, `when_all` captures any 1 exception

(*) for all policies currently in the Standard

Exceptions cause trouble

- **Parallel (C++17) algorithms**
 - `for_each`, `reduce`, `transform`, `sort`, ...
 - `ExecutionPolicy`: Permitted changes in execution order
 - Throw in loop body → terminate
- **Asynchronous tasks (C++11 `async`)**
 - Uncaught exception in task gets captured
 - Waiting on result throws passed-along exception
- **P0443 executors + P1897 asynchronous algorithms**
 - Separate path for handling ancestor task's uncaught exception
 - `when_all`: Express dependency on >1 tasks
 - If >1 parent throws, `when_all` captures any 1 exception

Exceptions cause trouble

- **Parallel (C++17) algorithms**
 - `for_each`, `reduce`, `transform`, `sort`, ...
 - `ExecutionPolicy`: Permitted changes in execution order
 - Throw in loop body → terminate
- **Asynchronous tasks (C++11 `async`)**
 - Uncaught exception in task gets captured
 - Waiting on result throws passed-along exception
- **P0443 executors + P1897 asynchronous algorithms**
 - Separate path for handling ancestor task's uncaught exception
 - `when_all`: Express dependency on >1 tasks
 - If >1 parent throws, `when_all` captures any 1 exception

Exceptions are for recoverable errors

Code characteristics leading to recoverable errors

Outside parallel code

Status of exceptions

May break user interface
if not carefully handled



Non-bug exceptions
& other recoverable
errors are more likely

Typical code activities

User interface, app driver

Launches &/or drives
parallel algorithms

Call 3rd-party libraries; do I/O;
complicated &/or speculative
computations

Inside parallel code

Exceptions usually
indicate bugs

Tight, optimized loops



Not allowed,
even if caught

Explicit SIMD; C++ dialects

Granularity (complexity of parallel “loop body” / task)

Example: Domain decomposition

- **Solve big linear system of equations $Ax=b$**
 - Decompose into many small systems (“subdomains”)
 - Solve small systems independently and combine results
 - Approximation → repeat to convergence
- **Domain decomposition might fail**
 - Entire problem has no solution (mathematically), or
 - Too many iterations to get an accurate solution, or
 - Some small systems may not have a solution
- **Fall-back algorithms take more memory and time**

Parallel domain decomposition

Previous non-parallel code

```
try {
    preSolveWork();
    for(auto&& subdomain: subdoms) {
        solve(subdomain);
    }
    postSolveWork();
} catch(/* what solve throws */) {
    useSlowerSolver = true;
}
```

Parallelization attempt

```
FixedMemoryPool pool(initSize);
try {
    preSolveWork();
    for_each(execution::par_unseq,
             begin(subdoms), end(subdoms),
             [&] (auto&& subdomain) {
                 solve(subdomain, pool);
             });
    postSolveWork();
} catch(/* what solve throws */) {
    useSlowerSolver = true;
}
```

Parallel domain decomposition

Previous non-parallel code

```
try {
    preSolveWork();
    for(auto&& subdomain: subdoms) {
        solve(subdomain);
    }
    postSolveWork();
} catch(/* what solve throws */) {
    useSlowerSolver = true;
}
```

- Fixed-size memory pool to speed up allocation in parallel loop
- Replace for loop with C++17 parallel algorithm

Parallelization attempt

```
FixedMemoryPool pool(initSize);
```

```
try {
    preSolveWork();
    for_each(execution::par_unseq,
             begin(subdoms), end(subdoms),
             [&] (auto&& subdomain) {
                 solve(subdomain, pool);
             });
    postSolveWork();
} catch(/* what solve throws */) {
    useSlowerSolver = true;
}
```

Parallel domain decomposition

Must distinguish 2 errors

- **Memory pool too small?**
 - Count total needed space
 - Reallocate pool
 - Retry domain decomposition
- **Any subdomain solve fails?**
 - Fall back to slower solver
- **Throw → terminate**
 - Can't distinguish by catching different exception types

Parallelization attempt

```
FixedMemoryPool pool(initSize);  
try {  
    preSolveWork();  
    for_each(execution::par_unseq,  
            begin(subdoms), end(subdoms),  
            [&] (auto&& subdomain) {  
                solve(subdomain, pool);  
            });  
    postSolveWork();  
} catch(/* what solve throws */) {  
    useSlowerSolver = true;  
}
```


Parallel domain decomposition

Must distinguish 2 errors

- **Memory pool too small?**
 - Count total needed space
 - Reallocate pool
 - Retry domain decomposition
- **Subdomain solve failed?**
 - Fall back to slower solver
- **Throw → terminate**
 - Can't distinguish by catching different exception types

Parallelization attempt

```
FixedMemoryPool pool(initSize);  
try {  
    preSolveWork();  
    for_each(execution::par_unseq,  
             begin(subdoms), end(subdoms),  
             [&] (auto&& subdomain) {  
                 solve(subdomain, pool);  
             });  
    postSolveWork();  
} catch(/* what solve throws */) {  
    useSlowerSolver = true;  
}
```

What do we do?!?

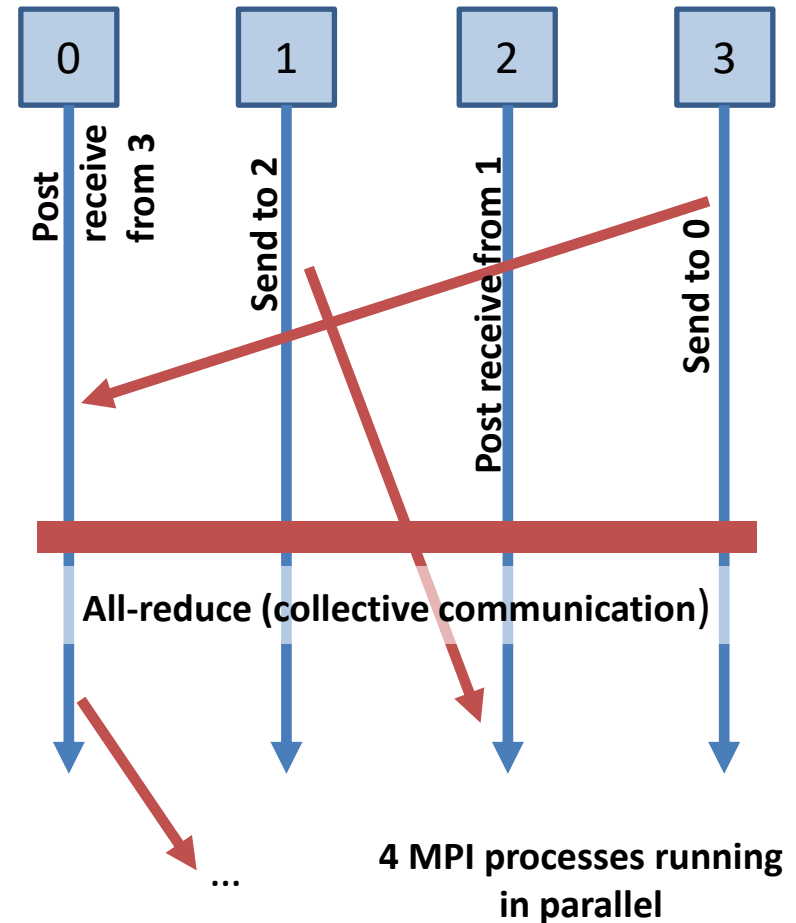
Message Passing Interface



- **C and Fortran interface**
 - For writing distributed-memory parallel code
- **Standard for 3 decades**
 - Unified divergent interfaces circa 1991
 - MPI 1.0 released 1994
 - 3.1 in 2015; 4.0 pending
- **Millions-way parallelism**
- **Stable: '90s code works**
- **Modest hardware requirements**
- **Cooperates with other programming models**
 - e.g., for GPUs

MPI's parallelism model

- **P “processes”**
 - Fixed location & number
 - Like `parallel_for_each` over 0, 1, ..., P-1
- **Do NOT share memory**
- **Communicate explicitly, mainly by messages**
 - Explicit function calls
 - “2-sided”: sender & receiver must participate
 - Point to point or collectives



MPI hostile to error handling

- ...a bit more so than C++ parallel algorithms
- **MPI deliberately punted on error handling**
 - Compared to early '90s competitor PVM
 - Goal: Run faster on wider range of systems
- **MPI equivalent of terminate: only best effort**
 - MPI_Abort on 1 process (hopefully) stops all processes
- **C++ exceptions: Undefined behavior if uncaught**
 - 1 process' uncaught exception often causes deadlock
 - Can set terminate_handler that calls MPI_Abort

MPI's error handling lessons

- **Turn exceptions into values**
 - before they risk breaking control flow or being uncaught
- **Turn fail-able “parallel for” into reduction**
 - Reduce on “did everybody succeed?”
 - Collect info for recovery and/or reporting
- **Prevent synchronization-related deadlock**
 - If you must synchronize...
 - ... use it as opportunity to communicate error state
- **Exploit out-of-band error reporting**

Turn exceptions to values, and reduce over work items

- **C++17 parallel algorithms**
 - Can't get uncaught exceptions back to caller, so:
 - Turn them to values (e.g., error code) via try-catch
 - Reduce over whether all work items succeeded
 - e.g., change `for_each` to `reduce`
 - See later for “not just a parallel for” algorithms like `sort`
- **P0443 tasks with P1897 `when_all`**
 - `when_all` drops all but 1 exception; if you need more:
 - Turn them to values via e.g., `let_error`
 - Next task “reduces” over values (like above)

Parallel domain decomposition

Reduce over error info

- **Each solve returns Result**
 - status: Why failed (bit field)
 - bytesNeeded (from pool)
- **Reduce over solve results**
 - If memory pool too small, this gives us required pool size

Parallelization

```
FixedMemoryPool pool(initSize);
preSolveWork();
Result result = transform_reduce(
    execution::par_unseq,
    begin(subdoms), end(subdoms),
    [] (Result x, Result y) {
        return Result{x.status & y.status,
            x.bytesNeeded + y.bytesNeeded}; },
    [&] (auto&& subd) {
        return trySolve(subd, pool);
    });
postSolveWork();
if(result.status & SOLVE_FAILED) {
    useSlowerSolver = true;
} else if(result.status & OUT_OF_MEM) {
    pool.resize(result.bytesNeeded);
    retry = true;
}
```

Parallel domain decomposition

Reduce over error info

- **Each solve returns Result**
 - status: Why failed (bit field)
 - bytesNeeded (from pool)
- **Reduce over solve results**
 - If memory pool too small, this gives us required pool size

Parallelization

```
FixedMemoryPool pool(initSize);
preSolveWork();
Result result = transform_reduce(
    execution::par_unseq,
    begin(subdoms), end(subdoms),
    [] (Result x, Result y) {
        return Result{x.status & y.status,
            x.bytesNeeded + y.bytesNeeded}; },
    [&] (auto&& subd) {
        return trySolve(subd, pool);
    });
postSolveWork();
if(result.status & SOLVE_FAILED) {
    useSlowerSolver = true;
} else if(result.status & OUT_OF_MEM) {
    pool.resize(result.bytesNeeded);
    retry = true;
}
```

Parallel domain decomposition

Reduce over error info

- **Each solve returns Result**
 - status: Why failed (bit field)
 - bytesNeeded (from pool)
- **Reduce over solve results**
 - If memory pool too small, this gives us required pool size

Parallelization

```
FixedMemoryPool pool(initSize);
preSolveWork();
Result result = transform_reduce(
    execution::par_unseq,
    begin(subdoms), end(subdoms),
    [] (Result x, Result y) {
        return Result{x.status & y.status,
            x.bytesNeeded + y.bytesNeeded}; },
    [&] (auto&& subd) {
        return trySolve(subd, pool);
    });
postSolveWork();
if(result.status & SOLVE_FAILED) {
    useSlowerSolver = true;
} else if(result.status & OUT_OF_MEM) {
    pool.resize(result.bytesNeeded);
    retry = true;
}
```

Parallel domain decomposition

Reduce over error info


- **Each solve returns Result**
 - status: Why failed (bit field)
 - bytesNeeded (from pool)
- **Reduce over solve results**
 - If memory pool too small, this gives us required pool size

Parallelization

```
FixedMemoryPool pool(initSize);
preSolveWork();
Result result = transform_reduce(
    execution::par_unseq,
    begin(subdoms), end(subdoms),
    [] (Result x, Result y) {
        return Result{x.status & y.status,
            x.bytesNeeded + y.bytesNeeded}; },
    [&] (auto&& subd) {
        return trySolve(subd, pool);
    });
postSolveWork();
if(result.status & SOLVE_FAILED) {
    useSlowerSolver = true;
} else if(result.status & OUT_OF_MEM) {
    pool.resize(result.bytesNeeded);
    retry = true;
}
```

Deadlock is easy in MPI


- Entire program is a “parallel for_each”
- Processes sometimes must communicate
- Each involved process must participate
- Else other processes may wait forever



Where a process may wait forever

Deadlock in parallel C++: Still possible

- **e.g., “thread barrier”**
 - All threads must reach it before any may proceed
 - e.g., `std::latch`
 - Can implement with atomic counter
- **Use: Shared resource**
 - Make sure all agents have stopped using it, so coordinator can release



Where an agent of execution may wait forever

Avoid deadlock by “going through the motions”

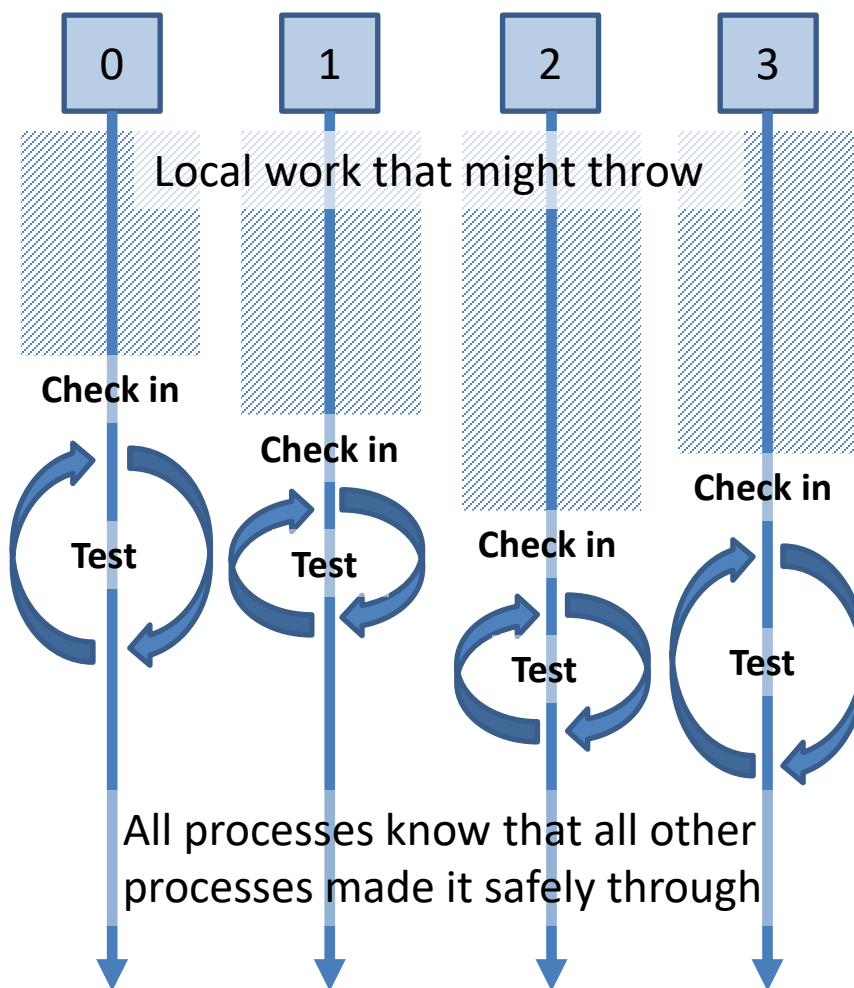
- **Think of parallel loop body as a sequence of “local” blocks, punctuated by synchronization**
 - Always participate in synchronization
 - Give each block a “bypass”: if error, do nothing harmlessly and pass along the running error state
 - Compare to P0443 “error channel”
- **Synchronization: opportunity for error reporting**
 - For thread barrier implemented as atomic counter:
 - Agent with error adds large number instead of one
 - (final count > number of agents) → error

Out-of-band error reporting

- **Asynchronous with respect to what you're using**
- **MPI: nonblocking barrier `MPI_Ibarrier`**
 - Collectives (barriers, reductions, ...) are not ordered with respect to other kinds of communication
 - Nonblocking != “makes progress in the background”
 - Communication need not happen until you wait or test
 - Analogous to `async(f, ...)`
 - But you can poll (test in a loop) to “force progress”
- **Use `MPI_Ibarrier` to test if a process threw**
 - Or dropped out

MPI: “Exception barrier”

- Run “risky” local work
 - Might throw
- “Check in” when done
 - Call `MPI_Ibarrier`
- Spin on `MPI_Test`
 - With a timeout
 - Can do other work speculatively while spinning
- Timeout → call `MPI_Abort`
- Success → safely through



Out-of-band in C++: Atomics

- **(Lock-free) atomic updates don't block**
 - ➔ can use them anywhere in parallel algorithms or tasks
 - e.g., write error info somewhere for later use
- **Use in sort or other non-reduction algorithms**
- **Use if you don't want to pay for reduction**
 - Not a zero-overhead abstraction if errors are rare
 - Atomic updates may hinder compiler optimizations
 - But that only matters if loop is “too optimized for recoverable exceptions” anyway

Summary

- **C++ parallel algorithms and asynchronous tasks make error handling harder**
- **Use design patterns developed for MPI**
 - Turn exceptions into values; reduce over error info
 - Avoid deadlock due to synchronization
 - Exploit out-of-band communication for error reporting
- **Thanks to Stellar Science!**

Summary

- **C++ parallel algorithms and asynchronous tasks make error handling harder**
- **Use design patterns developed for MPI**
 - Turn exceptions into values; reduce over error info
 - Avoid deadlock due to synchronization
 - Exploit out-of-band communication for error reporting
- **Thanks to Stellar Science!**

Hiring modern C++
software engineers!



hmeep hmeep