

The background of the slide is a dark, blue-toned photograph of a car, possibly a truck or SUV, driving on a bridge at night. The bridge's structural beams are visible, and the car's headlights are on. The overall mood is technical and modern.

# Apex.AI

Practical memory pool based  
allocators for Modern C++

by Misha Shalem  
[misha.shalem@apex.ai](mailto:misha.shalem@apex.ai)



- CppCon 2019:  
*Safe Software for Autonomous Mobility With Modern C++*  
by Andreas Pasternak
- Quote:  
*“Memory pools and allocators are only one piece of the solution”*

Today we going to talk about this one piece in (more) depth

# Memory allocations in real-time safety-critical environment

- The solution should be **safe** and **certifiable**
- What does it mean practically for C++ memory allocations?

We asked an independent 3<sup>rd</sup> party safety assessor and the answer was  
*“It should comply to Autosar C++ 14 Coding Guidelines regarding memory allocations”*

# Autosar C++ guidelines

## Rule A18-5-5 (required, implementation, partially automated)

Memory management functions shall ensure the following:

- *deterministic behavior resulting with the existence of worst-case execution time*
- *avoiding memory fragmentation*
- avoid running out of memory
- avoiding mismatched allocations or deallocations
- *no dependence on non-deterministic calls to kernel*


## Rule A18-5-7 (required, implementation, non-automated)

If non-realtime implementation of dynamic memory management functions is used in the project, then ***memory shall only be allocated and deallocated during non-realtime program phases***



# Memory pool


- Memory pools with fixed block sizes is a solution which addresses these issues:
  - It provides constant time (de)allocations in a preallocated buffer of memory
  - It avoids memory fragmentation with accurately-sized buckets
  - It is easy to implement, understand and reason about



# Memory pools and C++

There is a problem though...

The developer needs to know dimensions of all buckets in advance



$$\text{BucketSize} = \text{BlockSize} * \text{BlockCount}$$

# Memory pools and C++

- Maximum block count depends on the runtime
- **But block sizes should be known at compile time**

However...

C++ is all about high-level abstractions which hide the real sizes of allocations

Examples: `std::map`, `std::list`, `std::promise`, `std::shared_ptr`

All these can be parametrized with a user-provided allocator but they never actually use it directly

# Memory pools and C++

There is good news though

Even when the user-provided allocator is not used, a related (“rebound”) one usually is

And allocators are always familiar with the real types they allocate for

```
template<typename T>
class allocator {
    // ...
    template<typename U>
    struct rebound {
        using other = allocator<U>;
    };
    // ...
    using pointer = T*;
    // ...
    pointer allocate(size_type n, const void * hint = 0) {
        // has to allocate n * sizeof(T) bytes
    }
    // ...
};
```



# Goal

It should be possible to use any code which adopts standard C++ allocator model without modifications but allocate from a memory pool under the hood

# Implementation

## **Components:**

1. A memory pool
2. An allocator which allocates from a memory pool
3. A way to generate bucket definitions according to the actual allocations

# Memory pool implementation

## Components:

- 1. A memory pool***
2. An allocator which allocates from a memory pool
3. A way to generate bucket definitions according to the actual allocations

# Memory pool implementation

- A minimal implementation suffices as an example:
  - Some error checks are skipped for better readability
  - The implementation is not thread-safe
  - Default alignment is assumed
- However, multiple instances can be defined in the same application
  - Safety critical apps usually require separation of memory pools
  - A pool per thread would allow to avoid locking altogether

# Bucket implementation



# Bucket

```
class bucket {
public:
    const std::size_t BlockSize;
    const std::size_t BlockCount;

    bucket(std::size_t block_size, std::size_t block_count);
    ~bucket();
    // Tests if the pointer belongs to this bucket
    bool belongs(void * ptr) const noexcept;
    // Returns nullptr if failed
    [[nodiscard]] void * allocate(std::size_t bytes) noexcept;
    void deallocate(void * ptr, std::size_t bytes) noexcept;

private:
    // Finds n free contiguous blocks in the ledger and returns the first block's index or BlockCount on failure
    std::size_t find_contiguous_blocks(std::size_t n) const noexcept;
    // Marks n blocks in the ledger as "in-use" starting at 'index'
    void set_blocks_in_use(std::size_t index, std::size_t n) noexcept;
    // Marks n blocks in the ledger as "free" starting at 'index'
    void set_blocks_free(std::size_t index, std::size_t n) noexcept;

    // Actual memory for allocations
    std::byte* m_data{nullptr};
    // Reserves one bit per block to indicate whether it is in-use
    std::byte* m_ledger{nullptr};
};
```

# Construction/Destruction

```
bucket::bucket(std::size_t block_size, std::size_t block_count)
    : BlockSize{block_size}
    , BlockCount{block_count}
{
    const auto data_size = BlockSize * BlockCount;
    m_data = static_cast<std::byte*>(std::malloc(data_size));
    assert(m_data != nullptr);
    const auto ledger_size = 1 + ((BlockCount - 1) / 8);
    m_ledger = static_cast<std::byte*>(std::malloc(ledger_size));
    assert(m_ledger != nullptr);
    std::memset(m_data, 0, data_size);
    std::memset(m_ledger, 0, ledger_size);
}

bucket::~bucket() {
    std::free(m_ledger);
    std::free(m_data);
}
```

## bucket::allocate()

```
void * bucket::allocate(std::size_t bytes) noexcept {  
    // Calculate the required number of blocks  
    const auto n = 1 + ((bytes - 1) / BlockSize);  
  
    const auto index = find_contiguous_blocks(n);  
    if (index == BlockCount) {  
        return nullptr;  
    }  
  
    set_blocks_in_use(index, n);  
    return m_data + (index * BlockSize);  
}
```

## bucket::deallocate()

```
void bucket::deallocate(void * ptr, std::size_t bytes) noexcept {  
    const auto p = static_cast<const std::byte *>(ptr);  
    const std::size_t dist = static_cast<std::size_t>(p - m_data);  
    // Calculate block index from pointer distance  
    const auto index = dist / BlockSize;  
    // Calculate the required number of blocks  
    const auto n = 1 + ((bytes - 1) / BlockSize);  
    // Update the ledger  
    set_blocks_free(index, n);  
}
```

# Memory pool implementation



# How to configure the pools?

- A memory pool instance is merely a collection of buckets
  - Each bucket has two properties: **BlockSize** and **BlockCount**
  - A collection of these properties per bucket will define an instance of the pool
- 
- What would a collection of parameters look like?
  - How would we define more than one pool instance?

```
// The default implementation defines a pool with no buckets
template<std::size_t id>
struct bucket_descriptors {
    using type = std::tuple<>;
};
```

# Specializing instances

```
struct bucket_cfg16 {
    static constexpr std::size_t BlockSize = 16;
    static constexpr std::size_t BlockCount = 10000;
};

struct bucket_cfg32{
    static constexpr std::size_t BlockSize = 32;
    static constexpr std::size_t BlockCount = 10000;
};

struct bucket_cfg1024 {
    static constexpr std::size_t BlockSize = 1024;
    static constexpr std::size_t BlockCount = 50000;
};

template<>
struct bucket_descriptors<1> {
    using type = std::tuple<bucket_cfg16, bucket_cfg32, bucket_cfg1024>;
};
```

# Creating the memory pool(s)

```
template<std::size_t id>
using bucket_descriptors_t = typename bucket_descriptors<id>::type;

template<std::size_t id>
static constexpr std::size_t bucket_count = std::tuple_size<bucket_descriptors_t<id>>::value;

template<std::size_t id>
using pool_type = std::array<bucket, bucket_count<id>>;

template<std::size_t id, std::size_t Idx>
struct get_size
: std::integral_constant<std::size_t,
    std::tuple_element_t<Idx, bucket_descriptors_t<id>>::BlockSize>{};

template<std::size_t id, std::size_t Idx>
struct get_count
: std::integral_constant<std::size_t,
    std::tuple_element_t<Idx, bucket_descriptors_t<id>>::BlockCount>{};
```

# Creating the memory pool(s)

```
template<std::size_t id, std::size_t... Idx>
auto & get_instance(std::index_sequence<Idx...>) noexcept {
    static pool_type<id> instance{{{get_size<id, Idx>::value, get_count<id, Idx>::value} ...}};
    return instance;
}

template<std::size_t id>
auto & get_instance() noexcept {
    return get_instance<id>(std::make_index_sequence<bucket_count<id>>());
}
```

How to choose the right bucket?

How about simply going through the buckets until a big enough is found?

(Assuming they are sorted by block size )



# memory\_pool::allocate()

## Naive “First Fit” strategy

```
// Assuming buckets are sorted by their block sizes
template<std::size_t id>
[[nodiscard]] void * allocate(std::size_t bytes) {
    auto & pool = get_instance<id>();

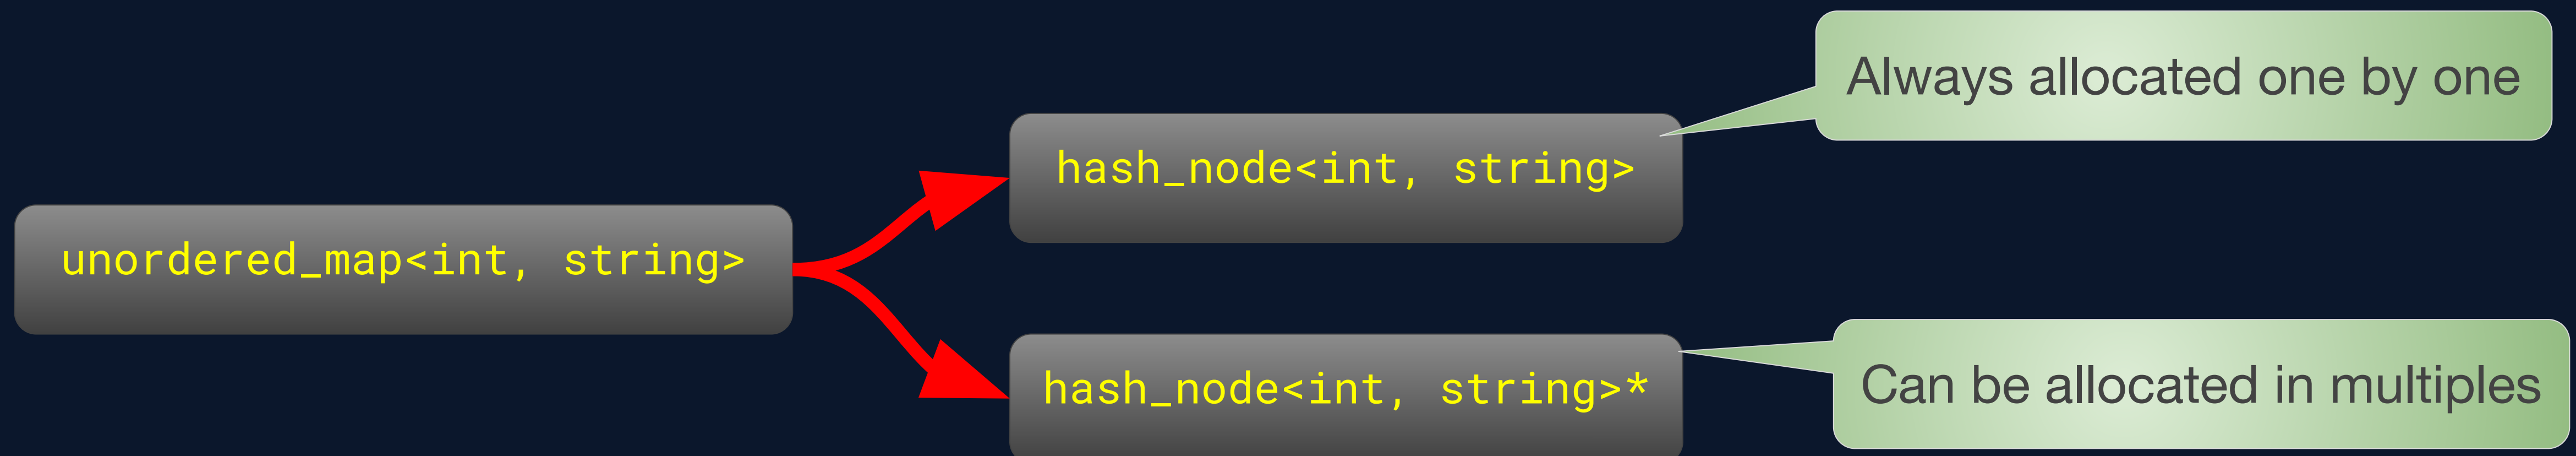
    for (auto & bucket : pool) {
        if(bucket.BlockSize >= bytes) {
            if(auto ptr = bucket.allocate(bytes); ptr != nullptr) {
                return ptr;
            }
        }
    }

    throw std::bad_alloc{};
}
```

## How to choose the right bucket?

This will work for the types which always allocate memory for their objects one by one (like `std::map`) but not for those which allocate memory for multiple objects at once (like `std::unordered_map`):

- There may simply be no bucket big enough
- A “wrong” bucket can be chosen which will be “stolen” from another type



```
memory_pool::allocate()
```

Closer to a real one

When allocating from a bucket it is unknown whether the allocation is for one or multiple objects — only the size in bytes is given

N.B. This also mimics how `std::pmr::memory_resource` defines the allocation interface

memory\_pool::allocate()

Closer to a real one

One way to solve this problem would be to find which allocation would lead to the least wasted memory and will take the least amount of blocks

```
struct info {
    std::size_t index{0}; // Which bucket?
    std::size_t block_count{0}; // How many blocks would the allocation take from the bucket?
    std::size_t waste{0}; // How much memory would be wasted?

    bool operator<(const info & other) const noexcept {
        return (waste == other.waste) ? block_count < other.block_count : waste < other.waste;
    }
};
```

# memory\_pool::allocate()

```
template<std::size_t id>
[[nodiscard]] void * allocate(std::size_t bytes) {
    auto & pool = get_instance<id>();
    std::array<info, bucket_count<id>> deltas;
    std::size_t index = 0;
    for (const auto & bucket : pool) {
        deltas[index].index = index;
        if (bucket.BlockSize >= bytes) {
            deltas[index].waste = bucket.BlockSize - bytes;
            deltas[index].block_count = 1;
        } else {
            const auto n = 1 + ((bytes - 1) / bucket.BlockSize);
            const auto storage_required = n * bucket.BlockSize;
            deltas[index].waste = storage_required - bytes;
            deltas[index].block_count = n;
        }
        ++index;
    }

    sort(deltas.begin(), deltas.end()); // std::sort() is allowed to allocate

    for (const auto & d : deltas)
        if (auto ptr = pool[d.index].allocate(bytes); ptr != nullptr)
            return ptr;

    throw std::bad_alloc{};
}
```



# How about fragmentation?

## Rule A18-5-5 (required, implementation, partially automated)

Memory management functions shall ensure the following:

- *deterministic behavior resulting with the existence of worst-case execution time*
- **avoiding memory fragmentation**
- *avoid running out of memory*
- *avoiding mismatched allocations or deallocations*
- *no dependence on non-deterministic calls to kernel*

Haven't we just reintroduced (local) memory fragmentation?

It depends on the logic of the users

They should be mindful of allocation/deallocation patterns and keep memory  
pools separated

# memory\_pool::deallocate()

```
template<std::size_t id>
void deallocate(void * ptr, std::size_t bytes) noexcept {
    auto & pool = get_instance<id>();

    for (auto & bucket : pool) {
        if (bucket.belongs(ptr)) {
            bucket.deallocate(ptr, bytes);
            return;
        }
    }
}
```

# Explicit initialization of the memory pool

```
template<size_t id> // Is there a specialization for this id?  
constexpr bool is_defined() noexcept {  
    return bucket_count<id>!= 0;  
}  
  
template<size_t id>  
bool initialize() noexcept {  
    (void) get_instance<id>();  
    return is_defined<id>();  
}
```

# Allocator implementation

## Components:

1. A memory pool
- 2. *An allocator which allocates from a memory pool***
3. A way to generate bucket definitions according to the actual allocations

# Allocator implementation

The allocator:

- Implements standard C++ allocator model compatible with `std::allocator_traits`
- Is bound to a specific memory pool instance in compile time
- Transparently allocates from another resource until memory pool is defined (*development*)
- Facilitates definitions of the buckets (*testing and production*)

# Usage scenarios

Using the same memory pool for nested objects:

```
static constexpr auto AllocId = 1;
using list = std::list<int, static_pool_allocator<int, AllocId>>;
using vector = std::vector<list, static_pool_allocator<list, AllocId>>;
// Both vector and list allocate from the pool with id 1
vector v;
v.emplace_back(5u, 42); // vector is [{ 42, 42, 42, 42, 42}]
```

Using separate memory pools for nested objects:

```
static constexpr auto ListAllocId = 1;
static constexpr auto VectorAllocId = 2;
using list = std::list<int, static_pool_allocator<int, ListAllocId>>;
using list_allocator = typename list::allocator_type;
using vector_allocator = static_pool_allocator<list, VectorAllocId>;
using scoped_adapter = std::scoped_allocator_adaptor<vector_allocator, list_allocator>;
using vector = std::vector<list, scoped_adapter>;
// List and vector allocate from different memory pools (1 and 2 respectively)
vector v;
v.emplace_back(5u, 42); // vector is [{ 42, 42, 42, 42, 42}]
```



# Static pool allocator

*The implementation skips the boilerplate code which is identical to*

`std::pmr::polymorphic_allocator`

# Static pool allocator

```
template<typename T = std::uint8_t, std::size_t id = 0>
class static_pool_allocator {
public:
    template<typename U>
    struct rebind { using other = static_pool_allocator<U, id>; };

    static_pool_allocator() noexcept : m_upstream_resource{pmr::get_default_resource()} {}
    static_pool_allocator(pmr::memory_resource * res) noexcept : m_upstream_resource{res} {}

    template<typename U>
    static_pool_allocator(const static_pool_allocator<U, id> & other) noexcept
        : m_upstream_resource{other.upstream_resource()} {}

    template<typename U>
    static_pool_allocator & operator=(const static_pool_allocator<U, id> & other) noexcept {
        m_upstream_resource = other.upstream_resource();
        return *this;
    }

    static bool initialize_memory_pool() noexcept { return memory_pool::initialize<id>(); }

private:
    pmr::memory_resource * m_upstream_resource;
};
```

# static\_pool\_allocator::allocate()

```
pointer allocate(size_type n, const void * = 0) {  
    // Reminder: here T is the type which we actually allocate for  
    if constexpr (memory_pool::is_defined<id>()) {  
        return static_cast<T*>(memory_pool::allocate<id>(sizeof(T) * n));  
    } else if (m_upstream_resource != nullptr) {  
        return static_cast<T*>(m_upstream_resource->allocate(sizeof(T) * n, alignof(T)));  
    } else {  
        throw std::bad_alloc{};  
    }  
}
```

# How to collect the type sizes?

- *Could we just log the sizes of  $T$  in the runtime?*
  - Does not necessarily reflect 100% line coverage
- *What if we were to make sure the line coverage is 100%?*
  - It could work but it can be impossible or unpractical to do for a real application
  - It would lack the call site information which would limit the insight users can get about where and why allocations happen, making the debugging and reasoning harder

It would be best to register all calls to `allocate()` instantiated for every type  $T$  in the compiled code, regardless of the runtime flow, as well as the call path which led to these allocations

# Injecting Instrumentation Code

```
namespace instrument {  
    template<std::size_t id, typename T, std::size_t size>  
    void type_reg() {}  
}
```

```
pointer allocate(size_type n, const void * = 0) {  
  
    instrument::type_reg<id, T, sizeof(T)>();  
  
    if constexpr (memory_pool::is_defined<id>()) {  
        return static_cast<T*>(memory_pool::allocate<id>(sizeof(T) * n));  
    } else if (m_upstream_resource != nullptr) {  
        return static_cast<T*>(m_upstream_resource->allocate(sizeof(T) * n, alignof(T)));  
    } else {  
        throw std::bad_alloc{};  
    }  
}
```

# static\_pool\_allocator::deallocate()

```
void deallocate(T * ptr, size_type n) {  
    if constexpr (memory_pool::is_defined<id>()) {  
        memory_pool::deallocate<id>(ptr, sizeof(T) * n);  
    } else if (m_upstream_resource != nullptr) {  
        m_upstream_resource->deallocate(ptr, sizeof(T) * n, alignof(T));  
    } else {  
        assert(false);  
    }  
}
```

# Allocation analyzer implementation

## Components:

1. A memory pool
2. An allocator which allocates from a memory pool
3. ***A way to generate bucket definitions according to the actual allocations***



# Allocation analyzer implementation

In order to collect the information injected by the allocator we need to:

- Compile the code with **clang** into LLVM bitcode
- Implement a custom LLVM pass(es) which would be able to:
  - Compose the list of all unique allocation types per pool instance
  - Record the call graph for every allocation (for debug purposes)
  - Generate bucket and memory pool definitions
- Run the pass on the bitcode using the LLVM **opt** tool

# Compile the code with clang into LLVM bitcode

- The code should be compiled with:
  - `-g` to have the *file:line* information for call graphs
  - `-O0` to prevent optimizing out our instrumentation calls
  - `-emit-llvm` to compile into the bitcode
  - Probably `-DNDEBUG` to get the *Release* versions of data structures

Then the bitcode files can be (optionally) linked together with `llvm-link`

# Printing the list of all unique allocations

```
class AllocListPass : public llvm::FunctionPass {
public:
    static char ID;
    AllocListPass() : llvm::FunctionPass(ID) {}

    bool runOnFunction(llvm::Function & f) override {
        const auto pretty_name = boost::core::demangle(f.getName().str().c_str());
        static const std::regex call_regex{R"(void instrument::type_reg<([^\, ]+), (.+), ([^\, ]+)>\(\))"};

        std::smatch match;
        if (std::regex_match(pretty_name, match, call_regex)) {
            if (match.size() == 4) {
                const auto pool_id = std::atoi(match[1].str().c_str());
                const auto type = match[2].str();
                const auto size = std::atoi(match[3].str().c_str());
                std::cout << "Pool ID: " << pool_id << ", Size: " << size << ", Type: " << type << "\n";
            }
        }
        return false; // does not alter the code, a read-only pass
    }
};

char AllocListPass::ID = 0;
static llvm::RegisterPass<AllocListPass> dummy("alloc-list", "This pass lists memory pool allocations");
```

# Analyzing call graph for every allocation

```
0 // file: /home/program.cpp
1 int main() {
2     f();
3     return 0;
4 }
5
6 void f() {
7     x();
8 }
9
10 void x() {
11     std::list<int, static_pool_allocator<int, 3>> lst;
12     lst.push_back(1);
13     // ...
14 }
```

- This time we use `llvm::ModulePass`
- Run a **depth-first search** starting from entry points (e.g. `main()`)
- Every time `type_reg<>` is discovered:
  - Register allocation type
  - Print the the call graph for allocation
- Check for recursions along the way and drop these branches
- Finally, the list of unique allocations can be translated into a header file with memory pool specializations

Call graph for: **Pool ID: 3, Size: 24, Type: `std::__1::__list_node<int, void*>`:**

```
1. static_pool_allocator<std::__1::__list_node<int, void*>, 3ul>::allocate(unsigned long, void const*) called at /usr/include/c++/v1/memory:1547
2. std::__1::allocator_traits<static_pool_allocator<std::__1::__list_node<int, void*>, 3ul>>::allocate(static_pool_allocator<std::__1::__list_node<int, void*>, 3ul>&, unsigned long) called at /usr/include/c++/v1/list:1079
3. std::__1::list<int, static_pool_allocator<int, 3ul>>::__allocate_node(static_pool_allocator<std::__1::__list_node<int, void*>, 3ul>&) called at /usr/include/c++/v1/list:1569
4. std::__1::list<int, static_pool_allocator<int, 3ul>>::push_back(int const&) called at /home/program.cpp:12
5. x() called at /home/program.cpp:7
6. f() called at /home/program.cpp:2
```

# Call graph caveats

The call graph might be incomplete because of:

- Virtual calls
- Calls through a function pointer
- Inline assembler

Both approaches can be combined by completing the list with allocations which are not in the call graph (“orphaned”) as a compromise solution

# Generating memory pool definitions

```
opt -load alloc-analyzer.so -alloc-analyze -gen-hdr my_defs.hpp -entry-point "main"< home/program.bc -o /dev/null
```

```
// file: my_defs.hpp
struct bucket_cfg24 {
    static constexpr std::size_t BlockSize = 24;
    static constexpr std::size_t BlockCount = 10000;
};

/* More buckets for memory pool 3 */

template<>
struct bucket_descriptors<3> {
    using type = std::tuple</* more buckets */ bucket_cfg24 /* more buckets */>;
};

/* Other buckets and memory pool specializations */
```

Include, recompile...

# Summary

- Memory allocations is a fundamental part of C++
- Memory allocations have to be deterministic to be usable in safety-critical applications
- A memory pool with fixed bucket sizes can be a solution but it is complex to use with the C++ allocator model
- An instrumenting allocator combined with an analyzing tool can be a certifiable solution for this problem





# Thank you!

[misha.shalem@apex.ai](mailto:misha.shalem@apex.ai)

© 2020 Apex.AI, Inc.